TEAM MEDVERSE

Gate 31 | DT and evidence bridge

Measured integration checkpoint · 20261103T132213_435523Z

What changed

The original PatientDigitalTwin class is connected through an observation-aware, per-session adapter. Source-filtered retrieval has explicit dense and lexical modes, with exact-ID lookup and source-linked evidence.

Engineering results

Checks: 42/42. HTTP checks: 11/11. Identical workflow outcomes: 432/432. Original files unchanged: 1118.

Dense backend status

BLOCKED_DEPENDENCY_OR_MODEL. No semantic retrieval accuracy is claimed. No silent BM25 fallback.

What the DT test establishes

Demonstration defaults are not observed data. New consent metadata cannot freshen older telemetry. Receiving-envelope metadata is not converted into unprovided patient measurements.

Language measurements

Known Korean/English question pairs are replayed with their original target labels. The diagrams show per-question ranks, not only aggregate recall. Corpus source dates are recorded metadata, not newly verified legal applicability.

What has NOT changed

Gate 30 action selection and commitment are preserved. The two C01/C02 progressive failures and capable-method ties remain. This checkpoint does not implement commit-versus-wait control or trained-agent inference.

Next dependency

Execute the dense path locally where the pinned model is available. The Gate 32 commitment-control design is frozen, but its controller and new scenario/reference files are not implemented or evaluated in this release.

Claim boundary

Local modeled engineering evidence. No clinical deployment, independent legal review, real authorization, general multilingual capability, or ITU conformance is established.

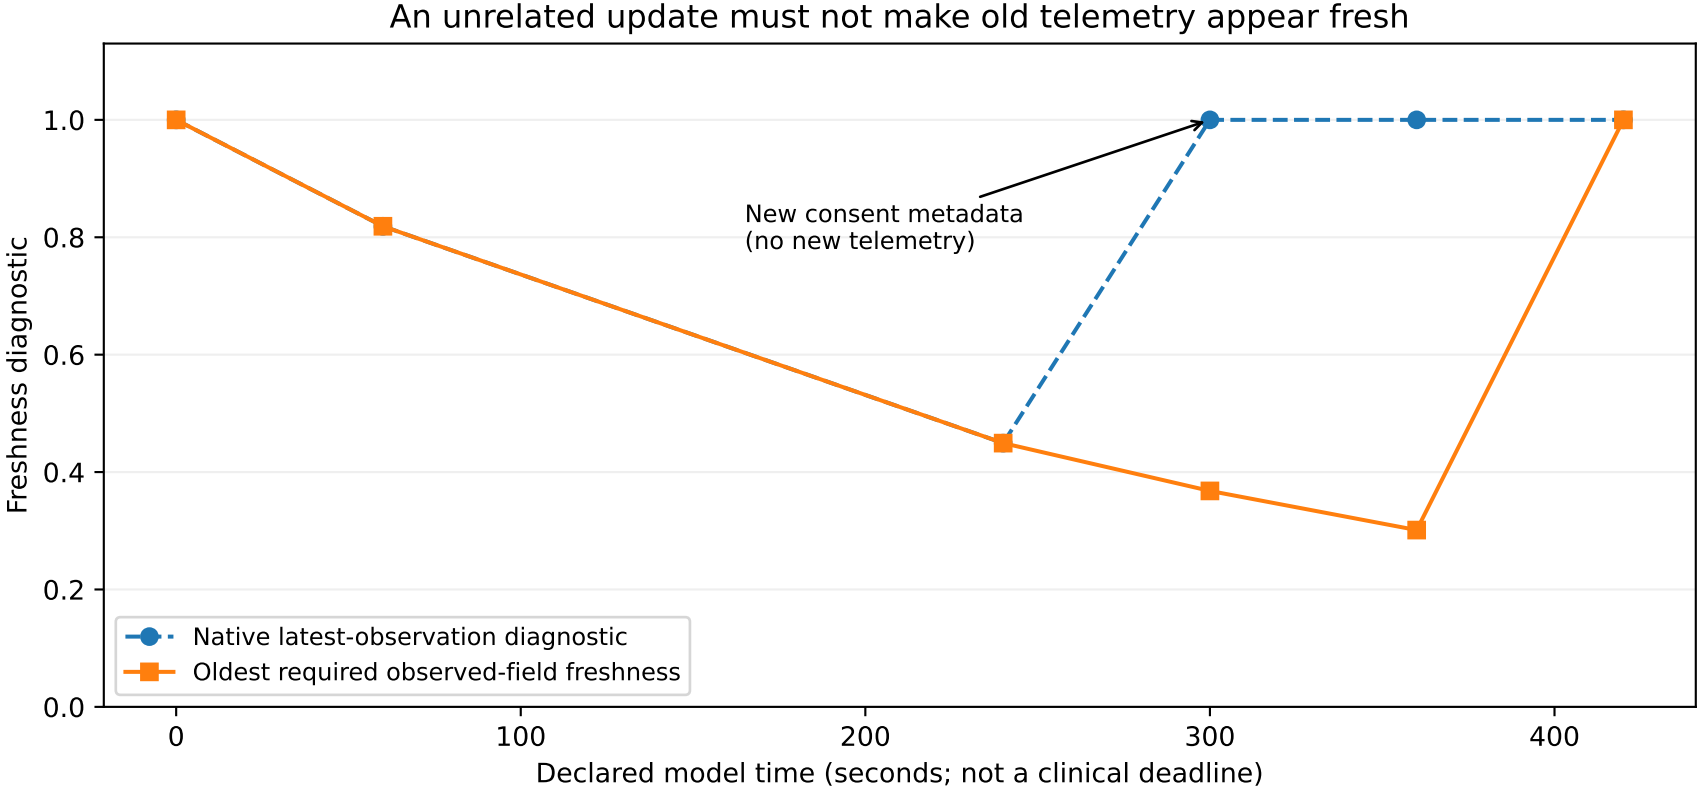

Explicit S1 development fixture. Native rules unchanged; adapter tracks per-field provenance.
No real patient observations, diagnosis, or action permission is established.

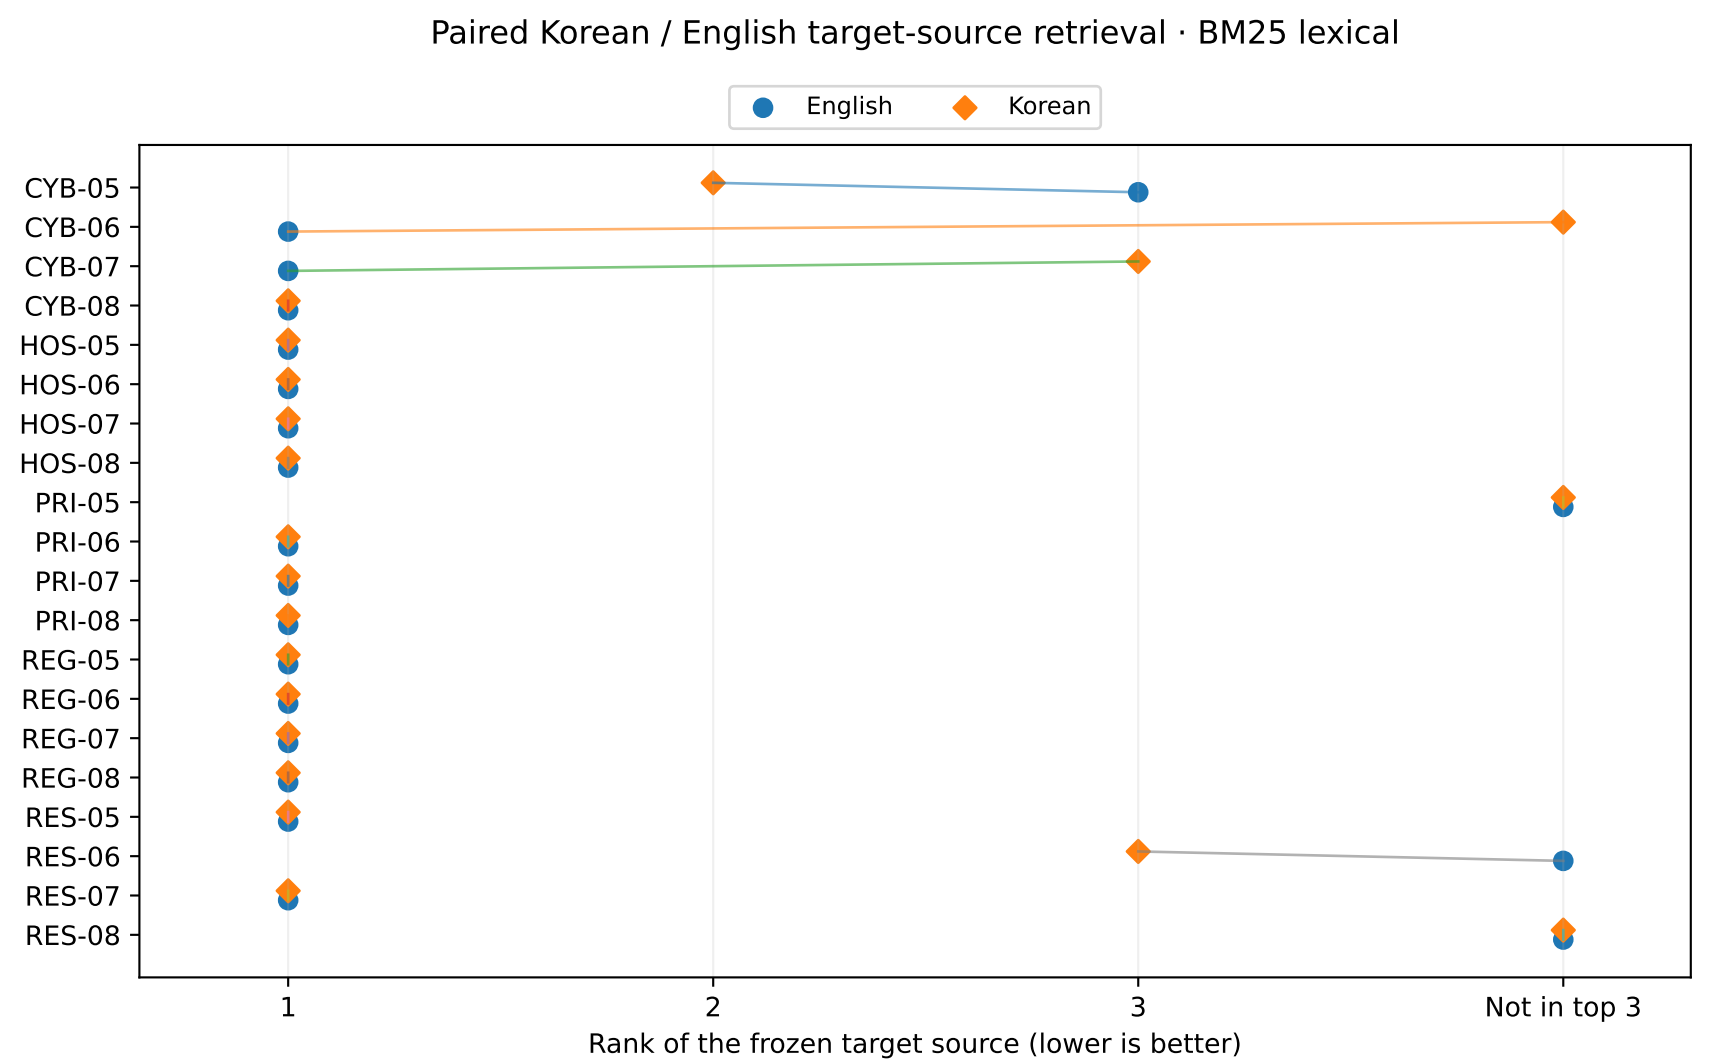

20 previously evaluated intent families, paired across languages; this is regression, not fresh confirmation.
Equal aggregate Recall@3 need not imply identical per-query success. Target labels were not revised.

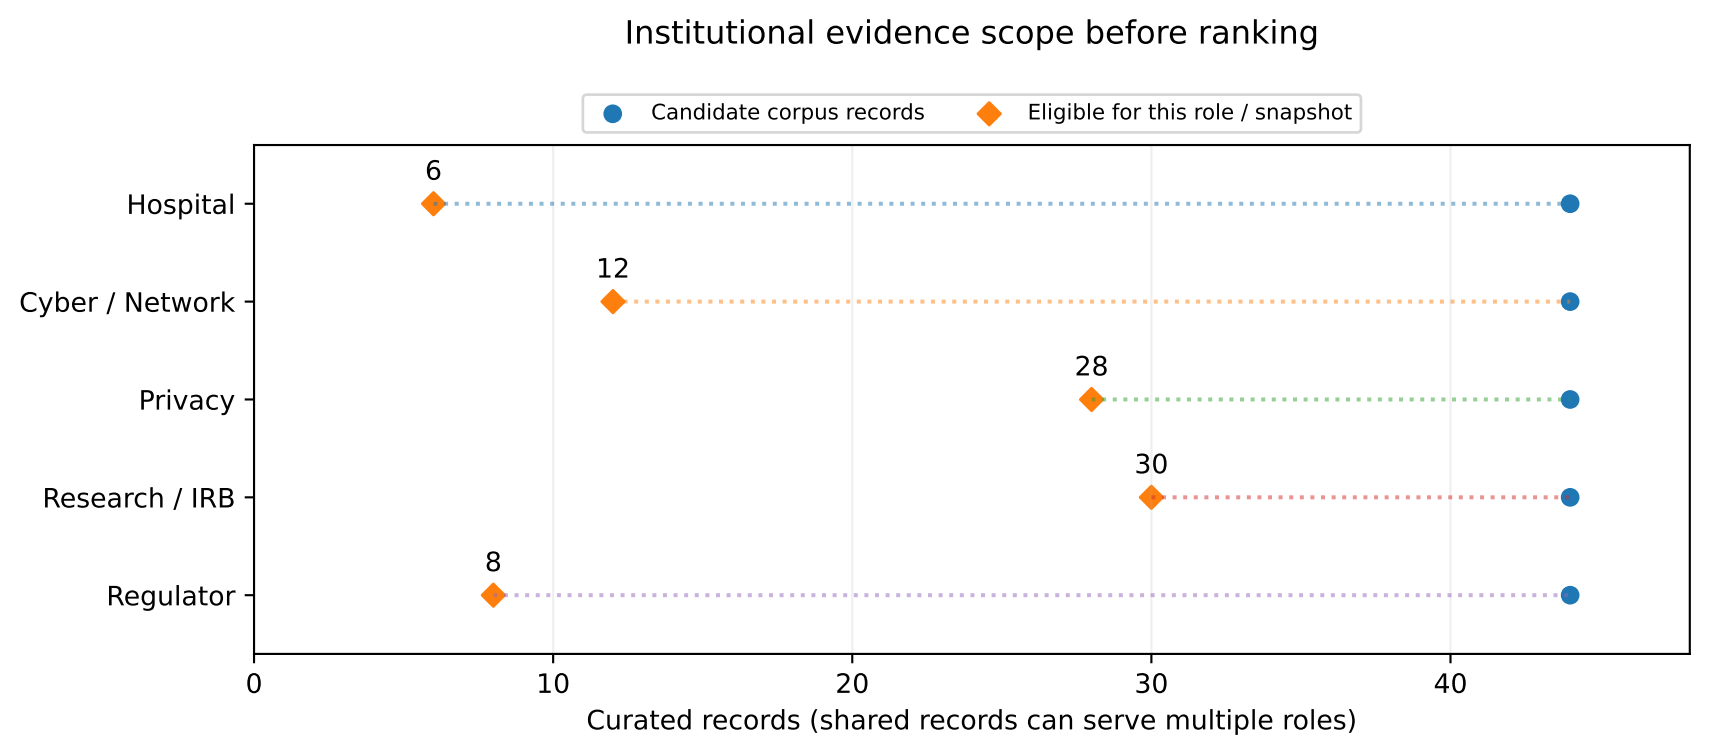

Counts measure configured retrieval scope, not completeness of Korean law or agent accuracy.
Policy applicability and executable interpretations still require review.